

Coursework 1: Randomized Decision Forest

Laura Palacio Garcia (CID: 01322823)
Imperial College London
lp2816@ic.ac.uk

Jenna Luchak (CID: 01429938)
Imperial College London
jkl17@ic.ac.uk

1. Training Decision Forest

Randomized decision forests generalize well to previously unseen data [2] and improve the prediction [5] thanks to randomizing the training dataset (Bagging) and the features (randomized node optimization). We performed Bagging (Bootstrap Aggregating) by randomly drawing samples of the same size from the training data [5], creating a new training subset for each tree. The samples were drawn with replacement so that the same data point could be chosen more than once. Figure 1 shows four different data subsets created, where some points might be the same within different subsets. Notice that the 3 different labels or classes correspond to the 3 main colors in the RGB model (red, green and blue).





Figure 1. Four data subsets generated by Bagging.

The size of each subset should be large because, otherwise, we would obtain a weak combined prediction. Our subsets have 95 samples, which corresponds to 63.2% of the total number of training data points ($N = 150$ samples). Alternatively, we created subsets of 150 samples with replacement, but the information gain was smaller. Larger subset sizes create less diversity among trees since there is almost no randomization (causing the ensemble model to have no effect). In this case, for large N , we would expect each subset to have 63.2% unique data points.

When splitting a node, the goal is to find the random split function that produces the highest information gain and confidence in the prediction, since we want the histograms of the child nodes to be more pure (exhibit a

value in only one class). In Figure 2 below, we tried two split functions: axis-aligned and two-pixel test (a type of linear function). There, we see that the distribution in the parent node is more or less uniform because we have almost the same number of points for each color (as in the right child node histograms). On the other hand, in Figure 2(b), the two-pixel test has a higher information gain since it has separated the data more successfully because the left child node only contains data points belonging to label 3 (color blue).



(a) Axis-aligned

(b) Two-pixel test

Figure 2. Two different split functions with the class histograms of the node and its two children nodes, and a measure of the information gains.

The information gain also depends on the degree of randomness ρ (*param.splitNum* in the code), which is the number of split functions we randomly evaluate at each node until choosing the best split function. Thus, if ρ is very small, every tree in the forest is very different between each other (high randomness). On the other hand, if ρ is large, then the algorithm finds the best split function after having tried a lot of random thresholds. Consequently, it is less random because all trees in the forest will be exactly the same (large tree correlation) [2].

After trying different split functions and degrees of randomness, we created Table 1 that shows the average information gain across all trees across 10 iterations. From Table 1, we see that the two-pixel test generates slightly larger information gain than axis-aligned and ρ does not seem to have a very significant effect (although it seems that $\rho = 5$ gives overall the best result). Given

the circular distribution of the data (as seen in Figure 2), neither the axis-aligned nor the two-pixel test are able to separate the data very well. Maybe a nonlinear weak learner could have worked better although it is computationally expensive and the other two functions have been proved to work well.

Weak learner	ρ	Information gain
Two-pixel test	3	0.1948
	5	0.1958
	10	0.1835
	15	0.1788
Vertical axis-aligned	3	0.1176
	5	0.1236
	10	0.1241
	15	0.1142
Horizontal axis-aligned	3	0.1186
	5	0.1233
	10	0.1184
	15	0.1165

Table 1. Information gain with respect to split function and degree of randomness used in randomised decision forest.

Finally, the tree has fully grown and, during testing, the input data reaches a leaf node, where it is classified using the class distribution at the leaf. In Figure 3, we show the class distributions for 9 leaf nodes. We can observe that 6 of them have a distribution such that, if an input data point reaches that leaf node, it will always be associated to a specific color (since that label has a probability of almost 1). However, the class or color prediction is harder in leaf nodes 5, 7 and 8.

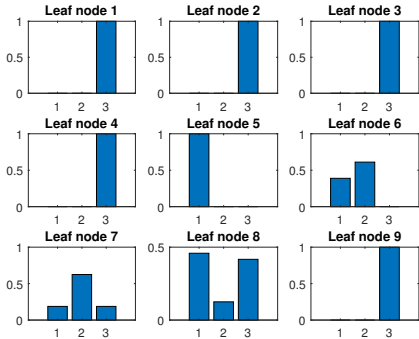


Figure 3. Class distributions of 9 leaf nodes.

The choice of stopping criteria influences the shape of the trees. Here, we stop the tree when a maximum number of depth has been reached because learning very deep trees increases the model complexity and can lead to overfitting. Also, we create a leaf node when the node contains less than 5 training points because it would not be useful to split the data more. Alternatively, we could have stopped when the information gain was smaller than a threshold.

2. Evaluating Decision Forest on the test data

In Figure 4, we show the class distributions of the leaf nodes (for each tree) reached by the third test point, as well as the averaged class distribution across all leaf nodes.

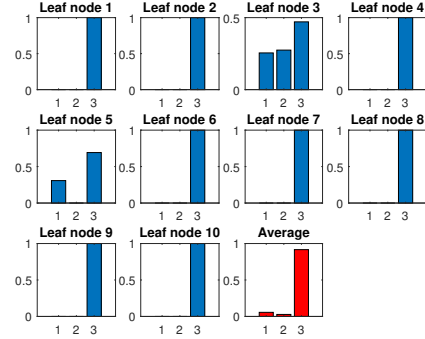


Figure 4. Class distributions of the leaf nodes the data point $[-0.7000, 0.4000]$ arrives to in every of the 10 trees and the averaged class distribution across all leaf nodes.

Then, we can show the predicted color for these four test points in Figure 5(a). The third test point is clearly blue since it had an average maximum probability of belonging to label 3 (as seen in Figure 4). As we observe in both subfigures in Figure 5, since the class distribution is a 3-dimensional vector of RGB values, we can plot the exact predicted color of each test point as a combination of the probabilities of the 3 main colors in RGB (instead of taking the label with the highest probability from the average class distribution).

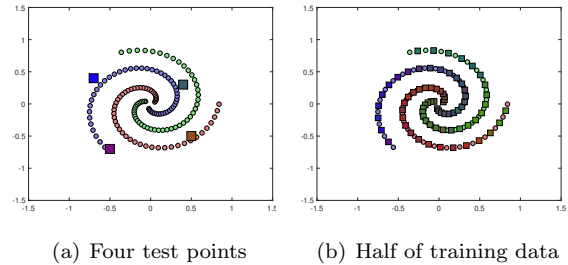


Figure 5. Color classification of different test points.

Our *data_test* from the Toy Spiral data set does not contain labels. Therefore, in order to test the accuracy of our model, we divided the training data again into two separate subsets: half of the data belongs to the training subset and the other half to the testing subset. This was done by saving alternate samples from each label to a new testing dataset and then erase those from the training dataset. Thus, this allows to have the same amount of data points for each color (label) in every subset. So, we will test our model with a subset

of data that was not used in the training process and the predicted colors are shown in Figure 5(b).

Then, we evaluated all data points in a dense 2D grid and the classification results obtained for different parameters are shown in Figure 6. The default parameters were 10 trees, maximum depth: 5, and $\rho = 3$.




Figure 6. Visualization of the classification results for all data points in a dense 2D grid for different values of number of trees (T) and tree depth (D).

As we can already see from Figure 6, we obtain more accurate predictions for larger number of trees, number of split functions (degree of randomness) and tree depth. When we increase the tree depth, there are more vertical lines in the dense grid so the prediction is not uniform in an area that should be the same color (as opposed to what happens when we increase the number of trees). With larger degree of randomness, the dense grid classification results also becomes smoother.

With the aim of observing this effect better, we calculated the accuracy of our model by using half the training data as testing data (as we did in Figure 5(b)). We take the true color of each data point and compare it with the color with the highest probability in the average class distribution across all leaf nodes. With the aim of generalizing our model, we repeated this process 10 times and obtained a 85.0667% average accuracy (using the default model parameters with an axis-aligned split function), whereas the two-pixel test obtained a better accuracy (89.7333%).

Then, we tested the classification performance with different values of number of trees (1, 3, 5, 10, 20, 30, 40, 50, 60, 70, 80, 90, 100), depth of trees (2, 5, 7, 11) and degree of randomness (2, 5, 7, 11, 15). We concluded that the testing accuracy increases with the number of trees, tree depth and degree of randomness (as in Figure 8). For instance, the accuracy ranged from 60% to 90% when we increased the number of

trees and tree depth. As for the values we tried for ρ , the accuracy increased from 82% to 88% with variations in between since it is not very stable.

Based on the figure we obtained, we decided that the optimal parameters for this classification problem are 40 trees with a maximum depth of 7, and a degree of randomness of 11. The main disadvantage of random forests is that it is hard to choose the value of these parameters. For instance, learning very deep trees increases the model complexity and can lead to overfitting, while the degree of randomness affects the generalization of the classification forest.

3. Experiment with Caltech101 dataset for image categorisation

3.1. K-means codebook

The vocabulary size of our K-means codebook is a matrix of size $[150 \times numBins]$. The number of rows in the codebook represents the number of images tested and trained in this data set; a constant value. There are 10 different classes, each has 15 images, for a total of 150. The number of columns is dependent on the variable $numBins$, which represents the number of clusters (bins) used in the algorithm K-means. The visual codebook is a set of cluster centers created by k-means that returns the centroid location of all clusters. The bags of visual words was constructed for every image using the codebook, via a vector quantization process that maps each patch vector to a certain code word.

When the default parameters for Matlab's built in function of K-means was executed, the simulation failed to converge. This behavior was expected, as K-means is time demanding and converges only to local optima; it depends on initialization. To solve this convergence issue, we implemented K-means using new techniques to reduce complexity. We tried increasing the maximum iterations of the algorithm, changing the initial clustering method from K-means++ algorithm (Matlab's default K-means algorithm) to the original K-means algorithm, employing an "Empty Action" parameter (*drop*) to remove any empty clusters (code words of frequency zero), and changing the number of clusters used; none of these techniques aided convergence. However, through this process we did learn that K-means++ computes slightly faster than K-means (139.081 seconds compared to 140.331 seconds), and a positive linear relationship exists between computation time and bin size. These two results were expected. When you look at the computational complexity equation of K-means: $O(DN'K)$, where N is the number of patches, K is the number of clusters and D is the descriptor dimensionality; complexity is

proportionally dependent on K . As number of clusters increases, computation time increases and complexity increases which causes K-means to perform more iterations every simulation. Maybe using Principal Component Analysis (which performs dimensionality reduction [3]) could have helped the algorithm converge.

The vector quantization process was an iterative process. Firstly, we initialize the bag-of-words matrix. Secondly, a three part *for* loop was initialized, to ensure that every patch from every image in every class of the data set was mapped to a code word. Thirdly, the distance between the i th patch vector and all of the code words was calculated, such that the minimum distance (using *vecnorm*) calculated was returned and the corresponding index of the code word was printed into a histogram vector. The bag-of-words histograms for the training and testing cat images are shown in Figure 7 below.

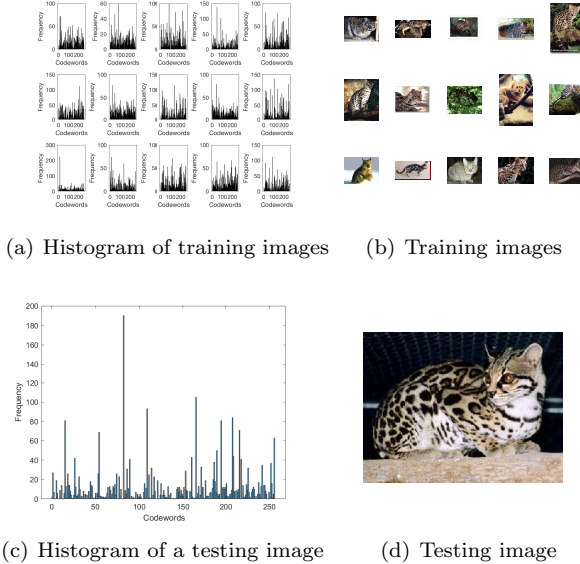


Figure 7. Bag-of-words histograms of example training and testing images from the category of wildcats.

We plotted histograms for all 15 images from class 7 (wildcat) for training and showed only one for a testing image for clarity. From Figure 7(a) and Figure 7(c), it was found that for a specific class all of the histograms expressed similar distributions. Most of the code words have an even frequency of patch vectors mapped to them, and a few unique outliers spiking to larger frequencies. This result suggests that our K-means codebook is fairly consistent at mapping the patch vectors of images in the same class. Although the histograms across the same class exhibit the same behavior, each image mapped a unique codebook, as can be seen by the varying frequencies across similar code words in Figure 7(a). This is expected, as wildcats have simi-

lar discernible features, such as color of fur, spots, and facial features but some of the images have distinctly different features such as the background colors or the fact that the cat is in a different position, a different size and less noticeable.

3.2. RF classifier

We used the training and testing data set, in the form of bag-of-words, obtained in Question 3.1, to train and test a random forest classifier. In order to determine the optimal classifier parameters, we investigated the vocabulary size, type of weak learner, number of trees, depth of trees, and degree of randomness (also called number of split functions ρ).

The approach we used to do the parameter sweep is cross-validation, which prevents over fitting. First, we randomly separate the training data (150 images) into K folds. In our case we chose $K = 5$ so that each subset would contain exactly 30 images. Then, the first 4 subsets (80% of the data) are used for training and the remaining one (20%) will be used for testing.. Therefore, we chose the parameters that produced a high average classification accuracy and a low standard deviation across the 5 folds and 10 iterations.

First, we explored visual vocabulary; how the number of code words affects the classification accuracy of our random forest classifier. The maximum number of colors an image can possibly display at one time is $2^8 = 256$ since image information is stored in 8-bit color graphics, while the lowest number of colors is 2^1 or 2 [1]. We chose to investigate vocabulary size by mapping our training data to a codebook of sizes 2^i , where $i = 1$ to 10. Random forest classifier has the same parameters for every trial: 40 numbers of trees, maximum depth of 8 trees, 20 split functions and an axis-aligned weak-learner. The reasoning behind the choice of these optimal parameters will be discussed throughout the investigation of this report.

At the same time, we investigated the type of weak-learner (axis-aligned and two-pixel test) used in the *splitNode* function of the random forest classifier. Therefore, all 10 vocabulary sizes were repeated for both weak-learners. The classification accuracy and total computation time for each size of codebook and each weak-learner can be seen in the Table 2 below.

From Table 2, we see that, in the case of axis-aligned weak learners, classification accuracy fluctuates up and down as the number of clusters used in K-means increases. The peak accuracy (60.67%) occurred at 256 clusters for the axis-aligned split function, whereas the accuracy of the two-pixel test never improved past 36.53%. The accuracy of the K-means algorithm, when using a proper weak learner, has a

Number of clusters	Axis-aligned		Two-pixel test	
	Accuracy [%]	Time [seconds]	Accuracy [%]	Time [seconds]
2	33.53	8.234	29.87	5.180
4	34.60	8.385	28.53	9.958
8	52.80	14.941	25.87	12.773
16	55.20	24.008	36.53	21.208
32	58.73	32.653	28.80	37.863
64	53.20	57.025	23.07	63.019
128	57.93	90.770	25.07	101.264
256	60.67	132.344	29.87	176.945
512	59.93	217.640	25.07	218.371
1024	46.93	259.450.315	22.53	258.968

Table 2. Classification accuracy and computation time with respect to the number of clusters and type of weak learner.

threshold codebook size since the accuracy does not increase for larger codebook sizes (1024) but the computation time extremely increases. Also, although each weak learner takes similar computation time for different sizes of codebooks, axis-aligned is comparatively more accurate than the two-pixel test across all codebook sizes. This result might be unexpected, as the two-pixel test is slightly more complex and supposed to perform better; although it always depends on the data available and what type of function is able to split it better, which in our case might be the axis-aligned.

The last parameters we investigated were number of trees, maximum depth of trees and ρ . We used axis-aligned and 256 clusters for this investigation, as per the conclusions from the previous parameter tests.




Figure 8. A random forest classifier when a specific parameter is manipulated and the average accuracy is calculated over 10 trials across 5 folds. Error bars are representative of the standard deviation across all trials.

The methodology for this investigation was iterative and involved three separate tests. For each test, one of the three parameters was manipulated while the other two parameters remained constant, this way we could observe the effects of each parameter individually. In the first iteration, constant parameters were used: 10 trees, tree depth: 5 and $\rho = 3$. However, for subsequent iterations, the constant parameters were changed to the optimal value obtained in the previous iteration.

The results from the first iteration are presented in Figure 8, where we see that, as number of trees increases from 0 to 70, the accuracy of the model increases and then plateaus. However, between 40 and 70 trees, accuracy does not improve by a significant margin (3%) and standard deviation is small. Since more trees increase complexity and computation time, 40 is optimal for this iteration. A similar relationship can be concluded for the manipulation of the maximum depth of trees and ρ ; the peak accuracies occurred at 8 and 20, respectively. The constant and optimal parameters chosen for the second iteration were number of trees 40, maximum tree depth of 8 and $\rho = 20$. This process was repeated until the optimal parameters for the current trial were equivalent to the previous trial. The parameters behaved similarly across all iterations with the exception of ρ in the final iteration because, as ρ changed, accuracy varied by at most 2%. Therefore, it can be concluded that the accuracy of the random forest classifier is dependent on all parameters; however, when number of trees and maximum depth of trees are optimal, ρ has no effect on accuracy. Also, the sensitivity of accuracy with respect to ρ was the least stable (largest standard deviation).

Using the optimal parameters found, we ran our simulation for 10 trials in order to calculate our model's optimal behavior, as well as build a confusion matrix. The total computation time for our model was 54.6440 seconds, with training time being 2.5922 seconds and testing 0.4505 seconds, which corresponds to 4.7438% and 0.8244% of the total computation time. Thus, we conclude that computation time was dominated by the K-means algorithm and training takes approximately 10 times longer than testing. The average accuracy and standard deviation were 60.2000% and 1.2192%, respectively.




Figure 9. Confusion Matrix of a random forest classifier that uses a K-means codebook mapped from 256 clusters.

The confusion matrix (Figure 9) is a 10 by 10 ma-

trix representing the correlation accuracies between the predicted and actual class labels given in the data set. The diagonal elements of the confusion matrix represent the percentage of instances that the predicted class was the actual class. We see in Figure 9 that the "trilobite", "wheelchair" and "wrench" classes were the most likely to be classified correctly (over 80%), while the class least likely to be classified correctly was the "umbrella" (20%).

Examples of some images that were correctly and incorrectly classified can be seen in Figure 10 below. A possible reason for misclassification may be because some images of different classes have similar features, such as the "water Lilly" and "umbrella" displaying similar colors and shapes, or the "yin and yang" symbol and "watch" both being round. Similar image features may lead to similarly mapped codebooks, thus misclassification.



(a) Correctly classified (b) Incorrectly classified

Figure 10. Example images randomly selected from the data set with the predicted labels.

3.3. RF codebook

Given the same data set, a new visual codebook was created using random forest. We grew trees with our descriptors matrix for all images (with the labels). Then, we tested the trees using the training matrix with all the descriptors and we obtained the leaves each specific descriptor for the image arrived at in each tree. Here, we used *testTrees_fast* instead of simply *testTrees* because it was significantly faster. Finally, the bag-of-words representations of the training image obtained by the RF codebook is obtained by adding the frequency of all the leaves for every image. We repeated the same process in order to create the testing images.

This training and testing data set, in the form of bag of words, was trained and tested using the same random forest classifier as detailed in Question 3.2. In Figure 11, we used again 5-fold cross-validation to compare the classification accuracy with respect to different random forest parameters (number of trees, maximum depth of trees, degree of randomness, weak learner) between the K-means codebook and the RF codebook.

If we compare the results from Figure 11 to those ob-

tained in Figure 8, we can see that the results are more or less the same if we compare building a RF codebook and a K-means one. When we used a two-pixel test as split function, the accuracy dropped to 29.1333%. Finally, there are two RF codebook parameters that can be changed: *PHOW_Sizes* and *PHOW_Step*, which correspond to the scales at which the dense SIFT features are extracted and the pixel steps of the grid at which the dense SIFT features are extracted (the lower, the denser), respectively [4]. No significant accuracy changes occurred after modifying these parameters.

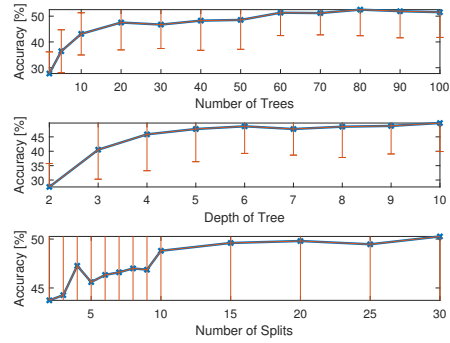


Figure 11. Same as in Figure 8 but using the RF codebook.

The computational complexity equation for K-means was outlined in Question 3.1, it is $O(DN'K)$, while the computational complexity equation for random forest is $O(\sqrt{DN'} \log K)$, where N is the number of patches, K is the number of clusters (or leaf nodes) and D is the descriptor dimensionality. Based on these formulas, it is clear to see that K and D have a larger affect on the complexity of K-means than in random forest. This statement is supported by the comparison in the computational time calculations (averaged across 10 trials) for both methods.

First, we found that building an RF codebook (28.7463 seconds) was much faster than a K-means codebook (54.6440 seconds). Then, we tested the optimal parameters found in Question 3.2 (40 trees, maximum depth: 8, degree of randomness: 20) and we obtained an average accuracy of 61.2667% with a standard deviation of only 1.7902%, while the average training time was 3.1123 seconds and the testing time was 0.5587 seconds. So, the accuracy is similar to when using the K-means codebook although the advantage of using an RF codebook is that it is much more computationally efficient since we have reduced the time to build a codebook to half. Since the data set in this project was large, the random forest was expected to perform more optimally and faster since labelling a descriptor with a balanced tree requires $O(\log K)$ operations whereas K-means costs $O(KD)$.

References

- [1] BBC Bitesize. Encoding images: Colour depth. [Online]. Available: <https://www.bbc.com/education/guides/zqyrq6f/revision/2>, 2018. [Accessed March 11, 2018].
- [2] Kari Pulli (Senior Director, NVIDIA Research). Machine Learning for Vision: Random Decision Forests and Deep Neural Networks. [Online]. Available: <http://web.stanford.edu/class/cs231m/spring-2014/docs/machine-learning.pdf>, 2014. [Accessed March 2, 2018].
- [3] J. Shlens. A tutorial on principal component analysis. *Measurement*, 51, 2005.
- [4] The VLFeat Authors. VLFeat - Documentation. [Online]. Available: http://www.vlfeat.org/matlab/vl_phow.html, 2007. [Accessed March 14, 2018].
- [5] Trevor Hastie (Stanford University). Trees, Bagging, Random Forests and Boosting. [Online]. Available: <http://jessica2.msri.org/attachments/10778/10778-boost.pdf>, 2004. [Accessed March 6, 2018].

Appendix

Question 1: Training Decision Forest

In order to show four data subsets obtained from Bagging, we added this code in function *growtrees.m* (inside the *for* loop after generating indexes for the training dataset):

```
1 % Plot subset training data used for
  this tree
2 if visualize
3     if T == 1
4         figure
5     end
6     if T <= 4
7         subplot(2,2,T)
8         plot_toydata(data(idx,:),);
9         title(['Subset ', num2str(T)])
10    end
11 end
```

The size of the subset is calculated in *growTrees.m*:

```
1 % Bootstrapping aggregating
2 % A new training set for each tree is
  generated by random sampling from
  dataset WITH replacement.
3 %idx = randsample(N,N,1); % same size
  as training data
4 idx = randsample(N,ceil(N*frac),1); %
  size is 63.2% of total training data
5 prior = histc(data(idx,red),labels)/
  length(idx);
```

Then, inside *splitNode.m*, we added the possibility to try another split function such as two-pixel test (using a *switch* statement that considers the split function we want to use or *mode*):

```
1 switch mode
2     case 'axis'
3         % Axis-aligned
4         m = 0; % input empty of visualise (
          used only for linear)
5         b = 0; % input empty of visualise (
          used only for linear)
6         dim = randi(D-1); % Pick one random
          dimension (1 is vertical and 2 is
          horizontal)
7         d_min = single(min(data(:,dim)))+eps;
          % Find the data range of this
          dimension
8         d_max = single(max(data(:,dim)))-eps;
9         t = d_min + rand*((d_max-d_min)); %
          Pick a random value within the
          range as threshold
```

```

10  idx_ = data(:,dim) < t;
11  case 'two'
12  % Two-pixel test
13  dim = 3; % used just to visualise (so
           it does not get confused with axis-
           aligned)
14  d_xmin = single(min(data(:,1))) + eps;
15  d_xmax = single(max(data(:,1))) - eps;
16  d_ymin = single(min(data(:,2))) + eps;
17  d_ymax = single(max(data(:,2))) - eps;
18  x_t = d_xmin + rand([1,2])*((d_xmax-
           d_xmin));
19  y_t = d_ymin + rand([1,2])*((d_ymax-
           d_ymin));
20  m = (y_t(2)-y_t(1))/(x_t(2)-x_t(1)); %
           slope
21  b = y_t(1) - m*x_t(1); % initial value
22  t = m*data(:,1) + b; % threshold
23  idx_ = data(:,2) < t;
24  end

```

To visualize the two-pixel test split function, we modified *visualise_splitfunc.m* so that it took it into account. So, we added the slope and the initial value of the function as inputs and we added another condition in the *if* statement:

```

1  function visualise_splitfunc(idx_best,
           data,dim,t,ig_best,iter,m,b) % Draw
           the split line
2  r = [-1.5 1.5]; % Data range
3
4  subplot(2,2,1);
5  if dim == 1 % vertical axis-aligned
6      plot([t t],[r(1),r(2)],'r');
7  elseif dim == 2 % horizontal axis-
           aligned
8      plot([r(1),r(2)],[t t],'r');
9  else % linear / two-pixel function
10     plot([r(1) r(2)],[m*r(1)+b m*r(2)+b
           ],'r');
11  end

```

The function to update the information gain also changes and we included new inputs and outputs:

```

1  function [node, ig_best, idx_best] =
           updateIG(node,ig_best,ig,t,idx,dim,
           idx_best,m,b) % Update information
           gain
2  if ig > ig_best
3      ig_best = ig;
4      node.t = t;
5      node.dim = dim;
6      idx_best = idx;

```

```

7      node.m = m; % update
8      node.b = b;
9  else
10     idx_best = idx_best;
11  end
12  end

```

In order to calculate the average best information gain for different split functions, we created a new output for *growTrees* called *global_ig*, which was a matrix containing the information gain for every tree for each time we split a node:

```

1  % Initialise base node
2  tree(T).node(1) = struct('idx',idx,'t',
           nan,'dim',-1,'prob',[],'m',[],'b'
           ,[]);
3
4  % Split Nodes
5  for n = 1:2^(param.depth-1)-1
6      [tree(T).node(n),tree(T).node(n*2),
           tree(T).node(n*2+1),ig_best] =
           splitNode(data,tree(T).node(n),
           param,mode);
7      global_ig(T,n) = ig_best;
8  end

```

Then, in *main_guideline*, we iterated *growTrees* 10 times and calculated the average information gain like this (*mode* is the type of split function we use):

```

1  % Grow all trees
2  for it = 1:10
3      [trees,global_ig] = growTrees(
           data_train,param,mode);
4      final_ig(it) = nanmean(nanmean(
           global_ig));
5  end
6  mean(final_ig)

```

Finally, instead of using function *visualise_leaf.m*, we preferred to create our own figure to visualize the first 9 leaf nodes and this is the code we added inside *main_guideline.m*:

```

1  figure
2  idx_tree = 1;
3  idx_node = 1;
4  for i = 1:9
5      try
6          prob_leaf = trees(idx_tree).leaf(
           idx_node).prob;
7      catch
8          idx_tree = idx_tree + 1;
9          idx_node = 1;

```



```

10     prob_leaf = trees(idx_tree).leaf(
        idx_node).prob;
11 end
12 subplot(3,3,i)
13 bar(prob_leaf);
14 title(['Leaf node ', num2str(i)])
15 idx_node = idx_node + 1;
16 end

```

Question 2: Evaluating Decision Forest on the test data

In this section, we tested the performance of our RF. This is the code written to Visualise the class distributions of the leaf nodes which the data point arrives at and the averaged class distribution:

```

1 all_probs = [];
2 for n = 1:size(test_point,1)
3     leaves = testTrees([test_point(n,:)
4         0],trees,mode);
5     % average the class distributions of
6     % leaf nodes of all trees
7     p_rf = trees(1).prob(leaves,:);
8
9     if length(trees) == 1
10         p_rf_sum = p_rf;
11     else
12         p_rf_sum = sum(p_rf)/length(trees);
13     end
14
15     % Plot class distributions of the
16     % leaf nodes
17     figure
18     for idx_tree = 1:length(trees)
19         subplot(3,4,idx_tree)
20         bar(p_rf(idx_tree,:));
21         title(['Leaf node ', num2str(
22             idx_tree)])
23     end
24     subplot(3,4,11)
25     bar(p_rf_sum, 'r')
26     title('Average')
27
28     all_probs(n,:) = p_rf_sum;
29 end

```

In order to evaluate all the data points in the dense 2D grid and visualise their classification results, we used this code:

```

1 %% Test on the dense 2D grid data, and
2 %% visualise the results
3 if test_train == 0
4     for n = 1:size(data_test,1)

```

```

4         leaves = testTrees(data_test(n,:),
5             trees);
6
7         % average the class distributions
8         % of leaf nodes of all trees
9         p_rf = trees(1).prob(leaves,:);
10        p_rf_sum = sum(p_rf)/length(trees);
11
12        all_probs_test(n,:) = p_rf_sum;
13    end
14    visualise(data_train, all_probs_test
15        ,0,0)
16 end

```

The *testTrees* function had to be changed so that it took into account the two-pixel test. So, the input *mode* specifies the type of split function that has been used for growing the trees.

```

1 t = tree(T).node(idx).t;
2 split_b = tree(T).node(idx).b;
3 split_m = tree(T).node(idx).m;
4 dim = tree(T).node(idx).dim;
5 % Decision
6 switch mode
7     case 'axis'
8         decision = data(m,dim) < t; % Pass
9         % data to left node (axis-aligned)
10    case 'two'
11        decision = data(m,2) < split_m*data(m
12            ,1) + split_b;
13    end
14
15    if decision % Pass data to left node (
16        % two-pixel test)
17        idx = idx*2;
18    else
19        idx = idx*2+1; % and to right
20    end

```

This is the code we created to separate the *data_train* equally into testing and training data:

```

1 % Separate training and testing
2 if test_train == 1
3     data_test = [];
4
5     % Label 1 (color red)
6     idx = find(data_train(:,3) == 1);
7     idx_test = idx(1:2:length(idx));
8     data_test = data_train(idx_test,:);
9     data_train(idx_test,:) = [];
10
11    % Label 2 (color green)
12    idx = find(data_train(:,3) == 2);

```

```

13 idx_test = idx(1:2:length(idx));
14 data_test = [data_test; data_train(
    idx_test,:)];
15 data_train(idx_test,:) = [];
16
17 % Label 3 (color blue)
18 idx = find(data_train(:,3) == 3);
19 idx_test = idx(1:2:length(idx));
20 data_test = [data_test; data_train(
    idx_test,:)];
21 data_train(idx_test,:) = [];
22 end

```

Then, we calculated the accuracy of our model like this:

```

1 %% Accuracy
2 if test_train == 1
3     true_labels = [];
4     predicted_labels = [];
5     for n = 1:size(test_point,1)
6         true_labels(n) = data_test(n,3);
7         [~, predicted_labels(n)] = max(
            all_probs(n,:));
8     end
9 end
10 accuracy = sum(true_labels ==
    predicted_labels)/size(test_point,1)
    *100;

```

Finally, we wanted to test the effect of different RF parameters on the accuracy of our model and we created three *for* loops to evaluate this. We used 5-fold cross-validation and, in order to provide proof that the model is generalized, we repeated the same process in 10 iterations. This is the loop created to test the effect of changing the number of trees and plot the average accuracy (the other two loops will be almost the same):

```

1 %% Change the RF parameter values and
    evaluate
2
3 % Set the random forest parameters
4 num_trees = [1, 3, 5, 10, 20, 30, 40,
    50, 60, 70, 80, 90, 100]; % Number
    of trees
5 depth_trees = [2, 5, 7, 11]; % Maximum
    depth of trees
6 splitNum_trees = [2, 5, 7, 11, 15]; %
    Number of split functions to try (p)
7
8 original_data_train = data_train;
9
10 % Reorder data randomly
11 randidx = randperm(size(
    original_data_train,1));

```

```

12 original_data_train =
    original_data_train(randidx,:);
13
14 % Cross-validation: Create the 5
    subsets
15 K = 5; % number of folds
16 subset_size = size(original_data_train
    ,1)/K; % 150/5 = 30 images for each
    subset
17
18 % 5-fold cross-validation: K-1 subsets
    for training and 1 for validation
19 for iter = 1:10
20     for fold = 1:K
21         % Validation subset (only one
            subset)
22         start_idx = subset_size*(fold
            -1) + 1;
23         stop_idx = start_idx +
            subset_size - 1;
24         data_test = original_data_train
            (start_idx:stop_idx,:); %
            train and test in same data
25
26         % Training subsets (the rest of
            the K-1 subsets)
27         data_train =
            original_data_train;
28         data_train(start_idx:stop_idx
            ,:) = []; % exclude the
            subset used for testing
29
30         %% Testing number of
            trees: param.num
31         for i = 1:length(param.num)
32             param.num = param_num(i);
33             param.depth = 8;
34             param.splitNum = 20;
35
36             % Train Random Forest
37             tic; % Start timer
38             tree = growTrees(data_train
                ,param,mode);
39             stop = toc; % Stop the
                timer
40             fprintf('TIC TOC Train
                random forest: %g\n',
                stop);
41
42             % Evaluate/Test Random
                Forest
43             tic; % Start timer
44             for n=1:size(data_test,1) %

```

```

45         Iterate through all
        rows of test data
        leaves = testTrees([
            data_test(n,:) 0],
            tree,mode); % Call
            the testTrees
            function
46         % average the class
            distributions of
            leaf nodes of all
            trees
47         p_rf = tree(1).prob(
            leaves,:);
48         p_rf_sum = sum(p_rf)/
            length(tree);
49         [~,predicted_label(n)]
            = max(p_rf_sum);
50     end
51     stop = toc; % Stop the
        timer
52     fprintf('TIC TOC Test
        random forest: %g\n',
        stop);
53
54     % Calculate accuracy of
        classifier
55     actual_label = data_test(:,
        end);
56     in_tree_accuracy(fold,i) =
        sum(actual_label ==
        predicted_label')/length
        (actual_label)*100;
57
58     end
59
60     % Calculate average and std
        accuracy across folds
61     trees_accuracy(iter,:) = mean(
        in_tree_accuracy);
62     trees_standard_dev(iter,:) = std(
        in_tree_accuracy);
63 end
64
65 % Calculate average and std accuracy
        across iterations
66 global_trees_accuracy = mean(
        trees_accuracy);
67 global_trees_standard_dev = mean(
        trees_standard_dev);
68
69 figure ,
70 subplot(3,1,1)
71 plot(param_num, global_trees_accuracy ,

```

```

        -x','linewidth',2);
72 hold on
73 errorbar(param_num,
        global_trees_accuracy ,
        global_trees_standard_dev);
74 xlim([min(param_num) max(param_num)]);
75 ylim([min(global_trees_accuracy) max(
        global_trees_accuracy)]);
76 xlabel('Number of Trees');
77 ylabel('Accuracy [%]');

```

Question 3: Experiment with Caltech101 dataset for image categorisation

Question 3.1: K-means codebook

The K-means codebook was obtained by modifying the function (getData). First, We built the clusters for each patch vector in the entire data set (*desc.tr*) by calling the algorithm for K-means from Matlab's toolbox. For our investigation on convergence, multiple calls of the K-means function were tested and are commented below the default. The number of clusters created by K-means is represented by variable "numBins" in our code. For our investigation on codebook size, this variable was manipulated. We also calculated the computation time it took the K-means function to cluster the data and printed it out. The code that we inserted into (getData) is as follows.

```

1 %% K-means clustering for each patch in
  the entire data set
2 % Call the kmeans function
3 % Inputs:
4 %   desc_sel': The transpose of the
  entire data set.
5 %
6 %   numBins: Number of clusters.
7 %
8 % Outputs:
9 %   index: 100,000 x 1 - Returns the
  cluster index for all
10 %   100,000 patch vectors
11 %
12 %   center: numBins x 128 returns the
  centroid location of each
13 %   cluster in desc_sel.
14 %
15 tic; % Initiate the timer
16 numBins = 256; % Number of clusters
  default.
17
18 % Call k means algorithm. Uncomment the
  line for investigation on
  convergence

```

```

19 [~,center] = kmeans(desc_sel',numBins,'
    MaxIter',100); % default
20 [~,center] = kmeans(desc_sel',numBins
    ', 'EmptyAction','drop');
21 [~,center] = kmeans(desc_sel',numBins
    ', 'Start','sample');
22
23 stop_kMeans = toc; % Stop the timer
24 fprintf('TIC TOC K-means: %g\n',
    stop_kMeans); % Print the time it
    takes to run kmeans

```

Second, we built a vector quantization method for the training data, such that every patch vector in every image was mapped to it's corresponding closest cluster. The data we added for vector quantization is as follows.

```

1 %% Vector Quantization training data
2 disp('Encoding Images...')
3
4 tic; % Initiate timer
5 numb_images = 0; % Initiate number of
    images counter
6 data_train = zeros(numel(desc_tr),
    numBins+1); % Initiate data_train
7
8 % For every single image in the
    training data set,
9 % determine the closest cluster from
    the kmeans algorithm relative
10 % to each patch vector, and create a
    Bag Of Words (BOW).
11 for r = 1:1:size(desc_tr,1) % Iterate
    through all rows
12     for c = 1:1:size(desc_tr,2) %
        Iterate through all columns
13         image = desc_tr{r,c}; % Define
            the current loops image
14         hist_vector = zeros(1,length(
            image)); % Reset
15
16         for i = 1:1:length(image) %
            Iterate through all patch
            vectors
17             db_I = double(image(:,i));
                % Convert to double
18             db_center = double(center')
                ; % Convert to double
                and transpose
19
20             % Calculate the distance
                between the current
                patch
21             % vector and all of the

```

```

    clusters
    distances = vecnorm(db_I -
        db_center);
22
23 % Determine the index of
    the cluster
    corresponding to
24 % the minimum distance and
    print it into a vector
    [~,cluster_num] = min(
        distances);
    hist_vector(i) =
        cluster_num;
25
26 end
27 % Use the histogram function to
    create a bag of words for
    every image
28 numb_images = numb_images + 1;
    % Update number of images
    counter
29 [BOW,~] = hist(hist_vector,
    numBins); % Generate bag of
    words vector
30
31 % Visual Vocabulary - Training
    Data
32 data_train(numb_images,1:end-1)
    = BOW; % Store BOW words
33 data_train(numb_images,end) = r
    ; % Label
34
35 if r == 7 %& (c==13 | c==14 | c
    ==15)
36     % Plot a histogram for
        training data bag of
        words
37     if c == 1
38         f1 = figure;
39         f2 = figure;
40     end
41     figure(f1)
42     subplot(3,5,c)
43     histogram(hist_vector,
        numBins);
44     xlabel('Codewords');
45     ylabel('Frequency');
46
47 % Plot the image
48 subFolderName = fullfile(
    folderName,classList{r});
49 imgList = dir(fullfile(
    subFolderName,'*.jpg'));
50 I = imread(fullfile(

```

```

        subFolderName,imgList(
            images_training(r,c)).
            name));
53         figure(f2)
54         subplot(3,5,c)
55         imshow(I);
56     end
57 end
58 end
59 stop = toc; % Stop the timer
60 fprintf('TIC TOC bag of words training:
        %g\n', stop);
61
62 % Clear unused variables to save memory
63 clearvars desc_tr desc_sel

```

The vector quantization process was replicated for the testing data, as follows.

```

1 %% Vector Quantization testing data
2 disp('Encoding Images...')
3
4 tic; % Initiate timer
5 numb_images = 0; % Initiate number of
   images counter
6 data_query = zeros(numel(desc_te),
   numBins+1); % Initiate data_test
7
8 % For every single image in the testing
   data set,
9 % determine the closest cluster from
   the kmeans algorithm relative
10 % to each patch vector, and create a
   Bag Of Words (BOW).
11 for r = 1:size(desc_te,1) % Iterate
   through all rows
12     for c = 1:size(desc_te,2) %
       Iterate through all columns
13         image = desc_te{r,c}; % Define
       the current loops image
14         hist_vector = zeros(1,length(
           image)); % Reset
15
16         for i = 1:length(image) %
           Iterate through all patch
           vectors
17             db_I = double(image(:,i));
               % Convert to double
18             db_center = double(center')
               ; % Convert to double
               and transpose
19
20             % Calculate the distance
               between the current

```

```

        patch
        % vector and all of the
        clusters
        distances = vecnorm(db_I -
            db_center);
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
        % Determine the index of
        the cluster
        corresponding to
        % the minimum distance and
        print it into a vector
        [~,cluster_num] = min(
            distances);
        hist_vector(i) =
            cluster_num;
        end
        % Use the histogram function to
        create a bag of words for
        every image
        numb_images = numb_images + 1;
        % Update number of images
        counter
        [BOW,~] = hist(hist_vector ,
            numBins); % Generate bag of
        words vector
        % Visual Vocabulary – Testing
        Data
        data_query(numb_images,1:end-1)
            = BOW; % Store BOW words
        data_query(numb_images,end) = r
            ; % Label
        if r == 7 %& (c==13 | c==14 | c
            ==15)
            % Plot a histogram for
            training data bag of
            words
            if c == 1
                f3 = figure('Units','
                    normalized','Position'
                    ,[.5 .1 .4 .9]);
                subtitle('Testing image
                    representations:
                    256-D histograms');
                f4 = figure;
            end
            figure(f3)
            subplot(3,5,c)
            histogram(hist_vector ,
                numBins);
            xlabel('Codewords');

```



```

49         ylabel('Frequency');
50
51     % Plot the image
52     subFolderName = fullfile(
53         folderName, classList{r});
54     imgList = dir(fullfile(
55         subFolderName, '*.jpg'));
56     I = imread(fullfile(
57         subFolderName, imgList(
58             images_testing(r,c)).name
59         ));
60     figure(f4)
61     subplot(3,5,c)
62     imshow(I);
63 end
64 end
65 stop = toc; % Stop the timer
66 fprintf('TIC TOC bag of words testing:
67         %g\n', stop);
68 % Clear unused variables to save memory
69 clearvars desc_te

```

Question 3.2: RF classifier

In order to build our random forest classifier, we first called the K-means codebook created in function *getData*. Please note that the entire RF classifier is inside of a *for* loop so that we were able to investigate the classifiers accuracy across different number of clusters. After obtaining the output variables from *getData* the RF classifier was built. For every cluster size tested, the classifier was run 10 times, and from those results, the average accuracy and standard deviation of those accuracies was returned. The code for these two steps are as follows.

```

1  init;
2  numBins = 256; % Optimal number of
3  clusters
4  %numBins = [2 4 8 16 32 64 128 256 512
5  1024]; % Uncomment for testing how
6  number of clusters affects k means
7  mode = 'axis';
8  for i=1:length(numBins)
9      %% Collect data from K means
10     Codebook
11     [data_train, data_test, stop_kMeans
12     (i), images_training,
13     images_testing] = getData('
14     Caltech', numBins(i));

```

```

9  %% Build RF classifier
10
11  for it = 1:10 % Iterate the random
12  forest classifier over 10 trials
13  % Set the random forest
14  parameters
15  param.num = 40; % The number of
16  trees
17  param.depth = 8; % Maximum
18  depth of the trees
19  param.splitNum = 20; % Number
20  of set of split parameters
21  theta i.e. p = 3
22  param.split = 'IG'; % Objective
23  function 'information gain'
24  Degree of randomness
25  parameter
26
27  % Train Random Forest
28  tic; % Start timer
29  tree = growTrees(data_train,
30  param, mode);
31  stop_train(i) = toc; % Stop the
32  timer
33  fprintf('TIC TOC Train random
34  forest: %g\n', stop_train(i)
35  );
36
37  % Evaluate/Test Random Forest
38  tic; % Start timer
39  for n=1:size(data_test,1) %
40  Iterate through all rows of
41  test data
42  leaves = testTrees([
43  data_test(n,:) ], tree,
44  mode); % Call the
45  testTrees function
46  % average the class
47  distributions of leaf
48  nodes of all trees
49  p_rf = tree(1).prob(leaves
50  ,:);
51  p_rf_sum = sum(p_rf)/length
52  (tree);
53  [~, predicted_label(n)] =
54  max(p_rf_sum);
55  end
56  stop_test(i) = toc; % Stop the
57  timer
58  fprintf('TIC TOC Test random
59  forest: %g\n', stop_test(i))
60  ;

```

```

36 % Calculate accuracy of
37 classifier
38 actual_label = data_test(:,end)
39 ;
40 accuracy(it) = sum(actual_label
41 == predicted_label')/length
42 (actual_label)*100;
43 end
44
45 avg_accuracy(i) = mean(accuracy');
46 % Calculate average accuracy
47 standard_dev(i) = std(accuracy'); %
48 Calculuate stadard deviations
49 across trials
50 end

```

Next, we created an if statement to follow the *for* loop that trained and tested the K-means code book. If there is more than 1 size of codebook being investigated, a graph can be plotted to visual the change in accuracy of the model. The code used for this is as follows.

```

1 % Plot a graph of the classification
2 accuracy for each trial only if
3 testing
4 % various numbers of clusters for k
5 means
6 if length(numBins)>1
7 figure ,
8 plot(numBins, avg_accuracy , '-ob');
9 set(gca, 'fontsize', 16);
10 xlabel('Number of Clusters');
11 ylabel('Classification Accuracy');
12 axis([min(numBins) max(numBins) 25 50])
13 ;
14 end

```

Next, another *if* statement was generated for the RF classifier. If only one codebook size is being tested, a confusion matrix will be calculated and displayed. The code for this section is as follows.

```

1 %% Calculate the confusion matrix
2 num_classes = 10;
3 confusion_matrix = zeros(num_classes ,
4 num_classes); % initialize
5 for test_class = 1:num_classes
6 idx_test = find(actual_label ==
7 test_class);
8 if isempty(idx_test) == 0 % there
9 is at least one value of this
10 class
11 for pred_class = 1:num_classes

```

```

8 idx_pred = find(
9 predicted_label(idx_test
10 ) == pred_class);
11 if isempty(idx_pred) == 0 %
12 there is at least one
13 value of this class
14 confusion_matrix(
15 test_class ,
16 pred_class) = length
17 (idx_pred)/length(
18 predicted_label(
19 idx_test));
20 end
21 end
22 end
23 end

```

```

24 % Plot confusion matrix with colormap
25 figure
26 ax1 = axes('Position',[0 0 1 1], '
27 Visible','off');
28 ax2 = axes('Position',[.2 .2 .7 .7]);
29 imagesc(confusion_matrix*100)
30 c = colorbar;
31 c.Label.String = 'Classification
32 frequency [%]';
33 %title('Confusion matrix for 4-way
34 classification')
35
36 % Create strings of the percentage in
37 confusion matrix
38 value_perc = num2str(confusion_matrix
39 (:)*100, '%0.2f');
40 value_perc = [value_perc, repmat('%',
41 numel(confusion_matrix), 1)];
42 value_perc = strtrim(cellstr(value_perc
43 )); % remove any space padding
44
45 % Find the value in the middle of
46 square to place text with accuracy
47 [x,y] = meshgrid(1:size(
48 confusion_matrix,1));
49 strings = text(x(:),y(:),value_perc(:),
50 'HorizontalAlignment','center');
51 mid_pos = mean(get(gca, 'CLim')); % find
52 middle value of the color range
53
54 % Change text colors
55 text_colors = repmat(confusion_matrix
56 (:)*100 < mid_pos, 1, 3);
57 set(strings, {'Color'}, num2cell(
58 text_colors, 2));

```

```

39 % Define labels and replace spaces with
    "newline" for better visualisation
40 labels = {'tick','trilobite','umbrella',
    'watch','water lilly','wheelchair'
    ,...
    'wild cat','window chair','
    wrench','yin yang'};
41
42 % Change axis
43 set(gca,'XTick',1:size(confusion_matrix
44 ,1),'XTickLabel',labels,...
45 'YTick',1:size(confusion_matrix,1),
    'YTickLabel',labels);
46 %ytickangle(90); % rotate 90 degrees
    the Y axis labels
47
48 % Add the labels of actual and
    predicted class
49 axes(ax1) % sets ax1 to current axes
50 text(0.09,0.55,{ 'ACTUAL'; 'CLASS'})
51 text(0.45,0.12, 'PREDICTED CLASS')

```

Lastly, two figures were plotted. One to visualize images that were correctly classified by our RF classifier, and the other to display images misclassified. The code for this section is as follows.

```

1 %% Plot example success/failures
2
3 folderName = 'Caltech_101/101
    _ObjectCategories';
4 classList = dir(folderName);
5 classList = {classList(3:end).name}; %
    10 classes
6
7 numImages = 4; % to plot
8 success = predicted_label' ==
    actual_label;
9
10 % Find index of images correctly
    classified and the ones
    misclassified
11 idx_success = find(success == 1);
12 idx_failure = find(success == 0);
13
14 % Randomly select some images to plot
15 rand_images_success = randsample(
    idx_success,numImages);
16 rand_images_failure = randsample(
    idx_failure,numImages);
17
18 % Create global list images used for
    testing
19 all_images_testing = [];

```

```

20 for row = 1:size(images_testing,1)
21     all_images_testing = horzcat(
        all_images_testing,
        images_testing(row,:));
22 end
23
24 % Plot success and failures
25 fig_success = figure;
26 fig_failure = figure;
27 for m = 1:length(rand_images_success)
28     % Success
29     label = predicted_label(
        rand_images_success(m));
30     num_image_folder =
        all_images_testing(
        rand_images_success(m));
31
32     % Plot the correctly classified
        image
33     subFolderName = fullfile(folderName
        ,classList{label});
34     imgList = dir(fullfile(
        subFolderName, '*.jpg'));
35     I = imread(fullfile(subFolderName,
        imgList(num_image_folder).name))
        ;
36     figure(fig_success)
37     subplot(2,2,m)
38     imshow(I);
39     title(labels(label))
40
41     % Failure
42     pred_label = predicted_label(
        rand_images_failure(m));
43     true_label = actual_label(
        rand_images_failure(m));
44     num_image_folder =
        all_images_testing(
        rand_images_failure(m));
45
46     % Plot the misclassified image
47     subFolderName = fullfile(folderName
        ,classList{true_label});
48     imgList = dir(fullfile(
        subFolderName, '*.jpg'));
49     I = imread(fullfile(subFolderName,
        imgList(num_image_folder).name))
        ;
50     figure(fig_failure)
51     subplot(2,2,m)
52     imshow(I);
53     title(labels(pred_label))
54 end

```

```

55 figure(fig_success)
56 subtitle('Correctly classified images')
57
58 figure(fig_failure)
59 subtitle('Misclassified images')

```

Question 3.3: RF codebook

We created a new function *getData_rf* to obtain the training and testing images. First, we obtain the matrix with the descriptors for the training images (*desc_tr*), by copying the same code as in *getData*. Then, we add the labels in the last column and create a visual vocabulary by applying RF and building a tree. The new code we added after getting *desc_tr* is:

```

1 % Add labels
2 for idx_row = 1:size(desc_tr,1)
3     for idx_col = 1:size(desc_tr,2)
4         desc_tr_labels{idx_row,idx_col} = [
5             desc_tr{idx_row,idx_col}, repmat(
6                 idx_row,128,1)];
7     end
8 end
9
10 % Build visual vocabulary (codebook)
11 for 'Bag-of-Words method'
12 desc_sel = single(vl_colsubset(cat(2,
13     desc_tr_labels{:}), 10e4)); %
14     Randomly select 100k SIFT
15     descriptors for clustering
16
17 param.num = 10; % The number of trees
18 param.depth = 5; % Maximum depth of the
19     trees
20 param.splitNum = 3; % Number of set of
21     split parameters theta i.e. p = 3
22 param.split = 'IG'; % Objective
23     function 'information gain' Degree
24     of randomness parameter
25 tree = growTrees(desc_sel', param, 'axis'
26     );

```

Then, this tree is used to find the leaves where each descriptor of every image arrives at. The next step is to calculate the frequency of each leaf node in the image (count the number of times the descriptors arrive at that leaf). This is how we create the training images. The code is the following:

```

1 % Training images
2 total_num_leaves = tree(size(tree,2)).
3     leaf(end).label;
4 count_image = 1;

```

```

4 for i = 1:size(desc_tr,1)
5     for j = 1:size(desc_tr,2)
6         leaves = testTrees_fast(desc_tr{i,j}
7             ', tree); % for each image
8
9         % Frequency histogram
10        % Count each leaf
11        for num_leaf = 1:total_num_leaves
12            freq_leaves(num_leaf) = length(
13                find(leaves == num_leaf));
14        end
15
16        % figure
17        % plot(freq_leaves)
18        % xlabel('Leaf node');
19        % ylabel('Frequency');
20
21        % From image to frequency of leaves
22        % across all trees
23        images_train(count_image,:) =
24            freq_leaves; % Store BOW words
25        images_train(count_image,end) = i;
26        % Label
27        count_image = count_image + 1;
28    end
29 end

```

Then, we create the testing images in the same way:

```

1 %% TESTING
2 cnt = 1;
3 % Load Images -> Description (Dense
4     SIFT)
5 for c = 1:length(classList)
6     subFolderName = fullfile(folderName,
7         classList{c});
8     imgList = dir(fullfile(subFolderName,
9         '*.jpg'));
10    imgIdx_te = imgIdx{c}(imgSel(1)+1:sum
11        (imgSel));
12
13    for i = 1:length(imgIdx_te)
14        I = imread(fullfile(subFolderName,
15            imgList(imgIdx_te(i)).name));
16        if size(I,3) == 3
17            I = rgb2gray(I);
18        end
19        [~, desc_te{c,i}] = vl_phow(single(
20            I), 'Sizes', PHOW_Sizes, 'Step',
21            PHOW_Step);
22    end
23 end
24 % Training images

```

```

19 total_num_leaves = tree(size(tree,2)).
    leaf(end).label;
20 count_image = 1;
21 for i = 1:size(desc_te,1)
22     for j = 1:size(desc_te,2)
23         leaves = testTrees_fast(desc_te{i,j}
            },tree); % for each image
24
25         % Frequency histogram
26         % Count each leaf
27         for num_leaf = 1:total_num_leaves
28             freq_leaves(num_leaf) = length(
                find(leaves == num_leaf));
29         end
30
31 %
32 %         figure
33 %         plot(freq_leaves)
34 %         xlabel('Leaf node');
35 %         ylabel('Frequency');
36
37 % From image to frequency of leaves
    across all trees
38 images_test(count_image,:) =
        freq_leaves; % Store BOW words
39 images_test(count_image,end) = i; %
        Label
40
41     count_image = count_image + 1;
42 end
43 end

```

Finally, in *main_guideline.m*, we train and test the classifier using the training and testing images obtained from *getData_rf*. We inspected the optimal parameters by applying 5-fold cross-validation again, as shown in Appendix 3.3. The code we show now is the one we used when we trained and tested the model using the optimal parameters across 10 iterations:

```

1 init
2
3 % Build codebook
4 [data_train, data_test] = getData_rf();
    % RF codebook
5
6 % Set the random forest parameters
7 param.num = 40;
8 param.depth = 8;
9 param.splitNum = 20;
10 param.split = 'IG'; % Objective
    function 'information gain' Degree of
        randomness parameter
11 mode = 'axis';

```

```

12 for iter = 1:10
13     % Train Random Forest
14     tic; % Start timer
15     tree = growTrees(data_train,param,
        mode);
16     stop_train(iter) = toc; % Stop the
        timer
17
18     % Evaluate/Test Random Forest
19     tic; % Start timer
20     for n=1:size(data_test,1) % Iterate
        through all rows of test data
21         leaves = testTrees([data_test(n
            ,:) 0],tree,mode); % Call
            the testTrees function
22         % average the class
            distributions of leaf nodes
            of all trees
23         p_rf = tree(1).prob(leaves,:);
24         p_rf_sum = sum(p_rf)/length(
            tree);
25         [~,predicted_label(n)] = max(
            p_rf_sum);
26     end
27     stop_test(iter) = toc; % Stop the
        timer
28
29     % Calculate accuracy of classifier
30     actual_label = data_test(:,end);
31     accuracy(iter) = sum(actual_label
        == predicted_label')/length(
            actual_label)*100;
32 end
33
34 % Calculate average and standard
    deviation accuracy and computation
    time
35 std_accuracy = std(accuracy')
36 avg_accuracy = mean(accuracy)
37 avg_stop_train = mean(stop_train)
38 avg_stop_test = mean(stop_test)
39 avg_time = avg_stop_train +
    avg_stop_test

```